

We've provided a synthetic data set consisting of property maintenance records attached. It's in CSV format. The nature of each column should be obvious from its name. Try not to spend more than 4 hours completing the task. Use whatever tools that you are comfortable with, though, if possible, Python is preferred. Here are the instructions:

- Draw a star schema data model for this data that would best facilitate quick reporting.
- Write a data pipeline to load the data into your star schema, using docker to host the database.
- It would be great to see some tests for your code.
- Zip up your code and share it with us, including some very short documentation explaining your design decisions.